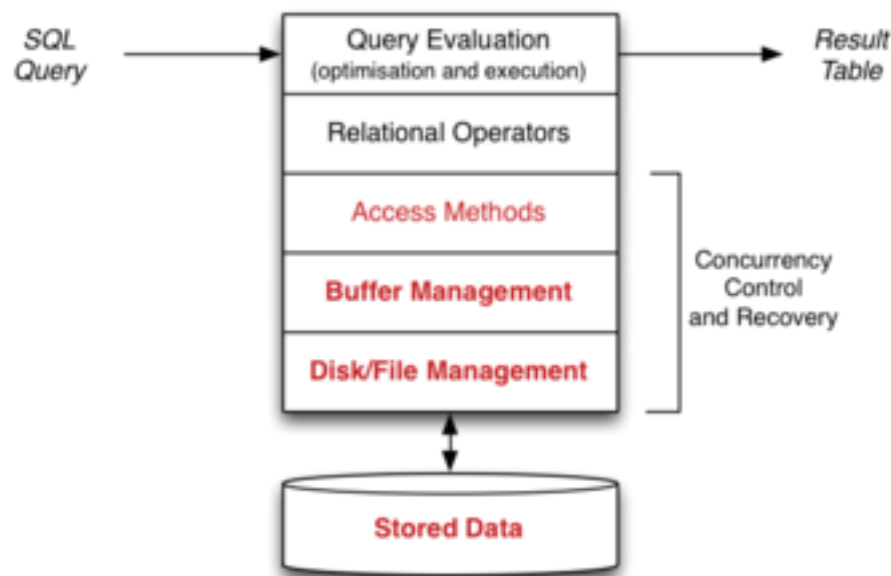


Storage Manager

Storage Management

2/100

Levels of DBMS related to storage management:



... Storage Management

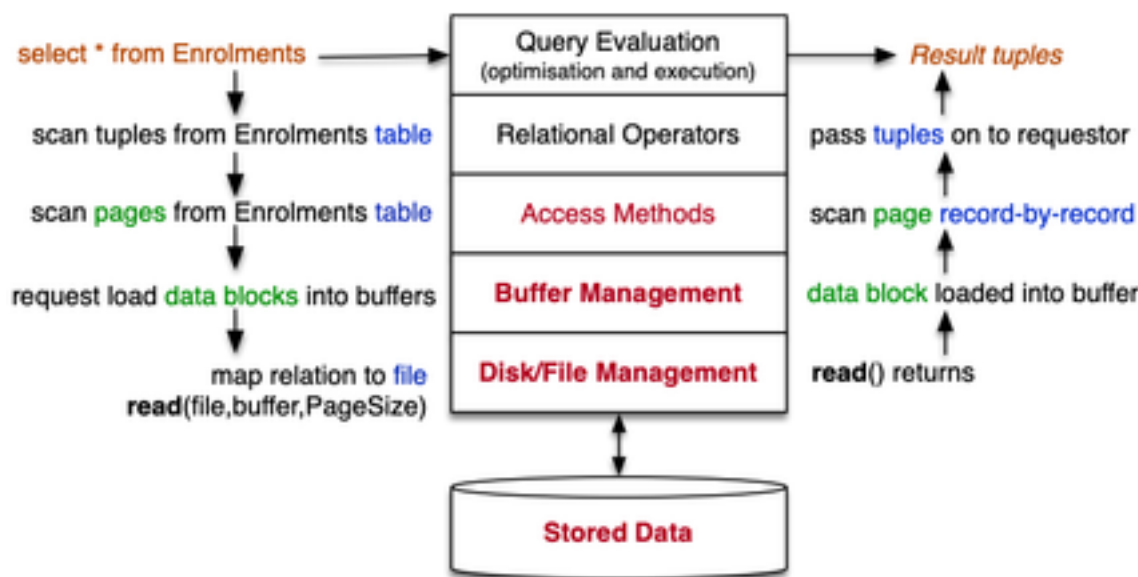
3/100

Aims of storage management in DBMS:

- map from database objects (e.g. tables) to disk files
- manage transfer of data to/from disk storage
- use buffers to minimise disk/memory transfers
- interpret loaded data as tuples/records
- provide view of data as collection of pages/tuples
- basis for "file structures" used by access methods

Views of Data in Query Evaluation

4/100



... Views of Data in Query Evaluation

5/100

Representing database objects during query execution:

- `DB` (handle on an authorised/opened database)
- `Rel` (handle on an opened relation)
- `Page` (memory buffer to hold contents of disk block)
- `Tuple` (memory holding data values from one tuple)

Addressing in DBMSs:

- `PageID = FileID+Offset` ... identifies a block of data
  - where `Offset` gives location of block within file
- `TupleID = PageID+Index` ... identifies a single tuple
  - where `Index` gives access to location of tuple within page

---

## Storage Management

6/100

Topics in storage management ...

- Disks and Files
  - performance issues and organisation of disk files
- Buffer Management
  - using caching to improve DBMS system throughput
- Tuple/Page Management
  - how tuples are represented within disk pages
- DB Object Management (Catalog)
  - how tables/views/functions/types, etc. are represented

---

## Storage Technology

### Storage Technology

8/100

Persistent storage is

- large, cheap, relatively slow, accessed in blocks
- used for long-term *storage* of data

Computational storage is

- small, expensive, fast, accessed by byte/word
- used for all *analysis/calculation* of data

Access cost HDD:RAM  $\cong$  100000:1, e.g.

- 10ms to read block containing two tuples
- $1\mu s$  to compare fields in two tuples

---

### ... Storage Technology

9/100

Hard disks are well-established, cheap, high-volume, ...

Alternative bulk storage: SSD

- faster than HDDs, no latency
- can read single items
- update requires block erase then write
- over time, writes "wear out" blocks
- require controllers that spread write load

Feasible for long-term, high-update environments?

Comparison of HDD and SSD properties:

	HDD	SDD
Cost/byte	~ 2c / GB	~ 13c / GB
Read latency	~ 10ms	~ 50μs
Write latency	~ 10ms	~ 900μs
Read unit	block (e.g. 1KB)	byte
Writing	write a block	write on empty block

Will SSDs ever replace HDDs for large-scale database storage?

Cost Models

Throughout this course, we compare costs of DB operations

Important aspects in determining cost:

- data is always transferred to/from disk as whole blocks (pages)
- cost of manipulating tuples in memory is negligible
- overall cost determined primarily by #data-blocks read/written

Complicating factors in determining costs:

- not all page accesses require disk access (buffer pool)
- tuples typically have variable size (tuples/page ?)

More details later ...

File Management

Aims of file management subsystem:

- organise layout of data within the filesystem
- handle mapping from database ID to file address
- transfer blocks of data between buffer pool and filesystem
- also attempts to handle file access error problems (retry)

Builds higher-level operations on top of OS file operations.

... File Management

Typical file operations provided by the operating system:

```
fd = open(fileName,mode)
// open a named file for reading/writing/appending
close(fd)
// close an open file, via its descriptor
nread = read(fd, buf, nbytes)
// attempt to read data from file into buffer
```

```
nwritten = write(fd, buf, nbytes)
// attempt to write data from buffer to file
lseek(fd, offset, seek_type)
// move file pointer to relative/absolute file offset
fsync(fd)
// flush contents of file buffers to disk
```

## DBMS File Organisation

14/100

How is data for DB objects arranged in the file system?

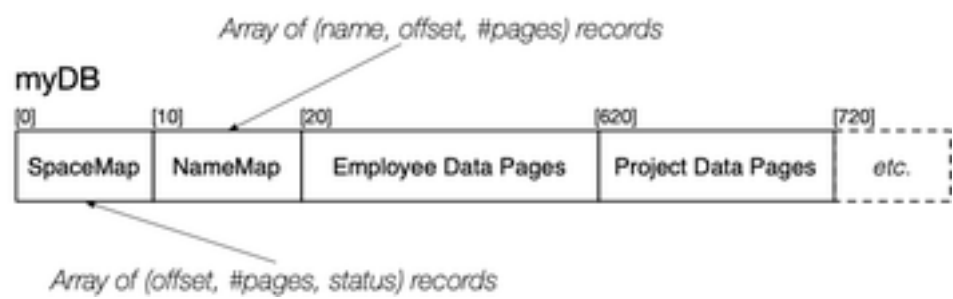
Different DBMSs make different choices, e.g.

- by-pass the file system and use a raw disk partition
- have a single very large file containing all DB data
- have several large files, with tables spread across them
- have multiple data files, one for each table
- have multiple files for each table
- etc.

## Single-file Storage Manager

15/100

Consider the following simple single-file DBMS layout:



E.g.

SpaceMap = [ (0,10,U), (10,10,U), (20,600,U), (620,100,U), (720,20,F) ]

NameMap = [ ("employee",20,350), ("project",620,40) ]

### ... Single-file Storage Manager

16/100

Each file segment consists of a number fixed-size blocks

The following data/constant definitions are useful

```
#define PAGE_SIZE 4096 // bytes per page

typedef long PageId; // PageId is block index
// pageOffset=PageId*PAGE_SIZE

typedef char *Page; // pointer to page/block buffer
```

Typical PAGE\_SIZE values: 1024, 2048, 4096, 8192

### ... Single-file Storage Manager

17/100

Storage Manager data structures for opened DBs & Tables

```
typedef struct DBrec {
    char *dbname; // copy of database name
    int fd; // the database file
    SpaceMap map; // map of free/used areas
    NameMap names; // map names to areas + sizes
```

```

} *DB;

typedef struct Relrec {
    char *relname; // copy of table name
    int start; // page index of start of table data
    int npages; // number of pages of table data
    ...
} *Rel;

```

---

## Example: Scanning a Relation

18/100

With the above disk manager, our example:

```
select name from Employee
```

might be implemented as something like

```

DB db = openDatabase("myDB");
Rel r = openRelation(db, "Employee");
Page buffer = malloc(PAGESIZE*sizeof(char));
for (int i = 0; i < r->npages; i++) {
    PageId pid = r->start+i;
    get_page(db, pid, buffer);
    for each tuple in buffer {
        get tuple data and extract name
        add (name) to result tuples
    }
}

```

---

## Single-File Storage Manager

19/100

```

// start using DB, buffer meta-data
DB openDatabase(char *name) {
    DB db = new(struct DBrec);
    db->dbname = strdup(name);
    db->fd = open(name, O_RDWR);
    db->map = readSpaceMap(db->fd);
    db->names = readNameMap(db->fd);
    return db;
}

// set up struct describing relation
Rel openRelation(DB db, char *rname) {
    Rel r = new(struct Relrec);
    r->relname = strdup(rname);
    // get relation data from map tables
    r->start = ...;
    r->npages = ...;
    return r;
}

```

---

## ... Single-File Storage Manager

20/100

```

// assume that Page = byte[PageSize]
// assume that PageId = block number in file

// read page from file into memory buffer
void get_page(DB db, PageId p, Page buf) {
    lseek(db->fd, p*PAGESIZE, SEEK_SET);
    read(db->fd, buf, PAGESIZE);
}

// write page from memory buffer to file
void put_page(Db db, PageId p, Page buf) {
    lseek(db->fd, p*PAGESIZE, SEEK_SET);
    write(db->fd, buf, PAGESIZE);
}

```

---

Consider a table  $R(x,y,z)$  with  $10^5$  tuples, implemented as

- number of tuples  $r = 10,000$
- average size of tuples  $R = 200$  bytes
- size of data pages  $B = 4096$  bytes
- time to read one data page  $T_r = 10msec$
- time to check one tuple  $1\ usec$
- time to form one result tuple  $1\ usec$
- time to write one result page  $T_r = 10msec$

Calculate the total time-cost for answering the query:

```
insert into S select * from R where x > 10;
```


if 50% of the tuples satisfy the condition.

DBMS Parameters

22/100

Our view of relations in DBMSs:

- a relation is a set of  $r$  tuples, with average size  $R$  bytes
- the tuples are stored in  $b$  data pages on disk
- each page has size  $B$  bytes and contains up to  $c$  tuples
- data is transferred disk↔memory in whole pages
- cost of disk↔memory transfer  $T_r, T_w$  dominates other costs



... DBMS Parameters

23/100

Typical DBMS/table parameter values:

Quantity	Symbol	E.g. Value
total # tuples	$r$	$10^6$
record size	$R$	128 bytes
total # pages	$b$	$10^5$
page size	$B$	8192 bytes
# tuples per page	$c$	60
page read/write time	$T_r, T_w$	10 msec
cost to process one page in memory	-	$\cong 0$

Multiple-file Disk Manager

24/100

Most DBMSs don't use a single large file for all data.

They typically provide:

- multiple files partitioned physically or logically
- mapping from DB-level objects to files (e.g. via meta-data)

Precise file structure varies between individual DBMSs.

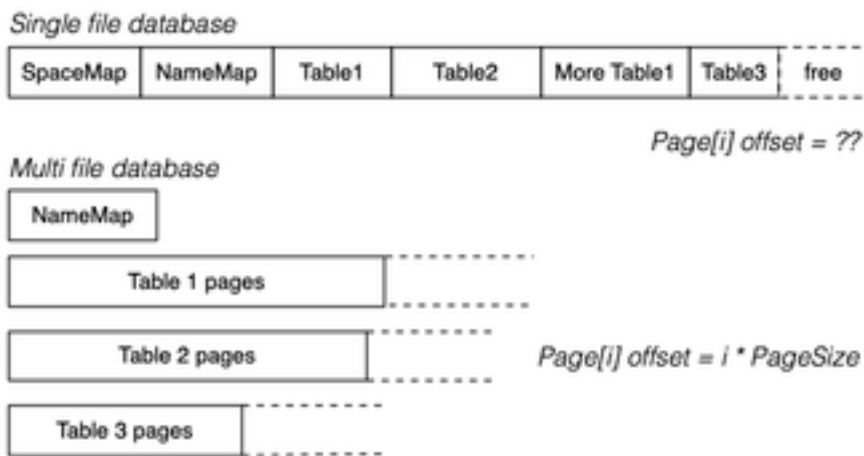
Using multiple files (one file per relation) can be easier, e.g.

- adding a new relation
- extending the size of a relation
- computing page offsets within a relation

## ... Multiple-file Disk Manager

25/100

Example of single-file vs multiple-file:



Consider how you would compute file offset of `page[i]` in `table[1]` ...

## ... Multiple-file Disk Manager

26/100

Structure of `PageId` for data pages in such systems ...

If system uses one file per table, `PageId` contains:

- relation identifier (which can be mapped to filename)
- page number (to identify page within the file)

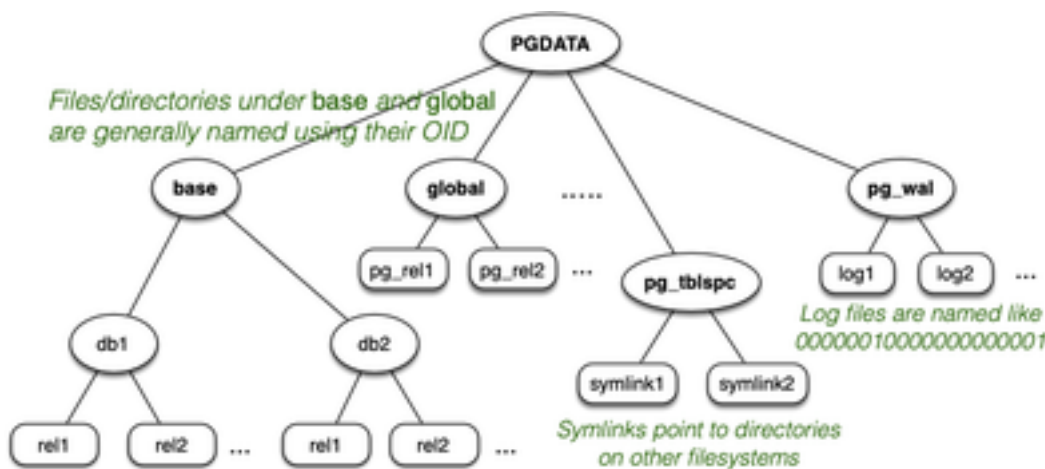
If system uses several files per table, `PageId` contains:

- relation identifier
- file identifier (combined with relid, gives filename)
- page number (to identify page within the file)

# PostgreSQL Storage Manager

27/100

PostgreSQL uses the following file organisation ...



## ... PostgreSQL Storage Manager

28/100

Components of storage subsystem:

- mapping from relations to files (**RelFileNode**)
- abstraction for open relation pool (**storage/smgr**)
- functions for managing files (**storage/smgr/md.c**)
- file-descriptor pool (**storage/file**)

PostgreSQL has two basic kinds of files:

- heap files containing data (tuples)
- index files containing index entries



## Relations as Files

PostgreSQL identifies relation files via their OIDs.

The core data structure for this is **RelFileNode**:

```
typedef struct RelFileNode {
    Oid   spcNode; // tablespace
    Oid   dbNode;  // database
    Oid   relNode; // relation
} RelFileNode;
```

Global (shared) tables (e.g. pg\_database) have

- spcNode == GLOBALTABLESPACE\_OID
- dbNode == 0

---

### ... Relations as Files

The **relpath** function maps **RelFileNode** to file:

```
char *relpath(RelFileNode r) // simplified
{
    char *path = malloc(ENOUGH_SPACE);

    if (r.spcNode == GLOBALTABLESPACE_OID) {
        /* Shared system relations live in PGDATA/global */
        Assert(r.dbNode == 0);
        sprintf(path, "%s/global/%u",
                DataDir, r.relNode);
    }
    else if (r.spcNode == DEFAULTTABLESPACE_OID) {
        /* The default tablespace is PGDATA/base */
        sprintf(path, "%s/base/%u/%u",
                DataDir, r.dbNode, r.relNode);
    }
    else {
        /* All other tablespaces accessed via symlinks */
        sprintf(path, "%s/pg_tblspc/%u/%u/%u", DataDir
                r.spcNode, r.dbNode, r.relNode);
    }
    return path;
}
```

---

## Exercise 2: PostgreSQL Files

In your PostgreSQL server

- examine the content of the \$PGDATA directory
- find the directory containing the pizza database
- find the file in this directory for the People table
- examine the contents of the People file
- what are the other files in the directory?
- are there *forks* in any of your databases?

---

## File Descriptor Pool

Unix has limits on the number of concurrently open files.

PostgreSQL maintains a pool of open file descriptors:

- to hide this limitation from higher level functions
- to minimise expensive open ( ) operations

File names are simply strings: **typedef char \*FileName**

Open files are referenced via: **typedef int File**

A **File** is an index into a table of "virtual file descriptors".

Defs: **include/storage/fd.h**  
 Code: **backend/storage/file/fd.c**



Interface to file descriptor (pool):

```
File PathNameOpenFile(char *fileName, int flags)
    // open a file with default pg.conf mode
File PathNameOpenFilePerm(char *fName, int flags, int mode)
    // open a file in the DB directory ($PGDATA/base/...)
File OpenTemporaryFile(bool interXact)
    // open temp file flag: close at end of transaction?
void FileClose(File file)
void FileUnlink(File file)
int FileRead(File file, char *buffer, int amount)
int FileWrite(File file, char *buffer, int amount)
int FileSync(File file)
long FileSeek(File file, long offset, int whence)
int FileTruncate(File file, long offset)
```

Analogous to Unix syscalls `open()`, `close()`, `read()`, ...

Virtual file descriptor records (simplified):


```
typedef struct vfd
{
    s_short fd;           // current FD, or VFD_CLOSED if none
    u_short fdstate;      // bitflags for Vfd's state
    File nextFree;        // link to next free Vfd, if in freelist
    File lruMoreRecently; // doubly linked recency-of-use list
    File lruLessRecently;
    long seekPos;         // current logical file position
    char *fileName;       // name of file, or NULL for unused Vfd
    // NB: fileName is malloc'd, and must be free'd when closing the Vfd
    int fileFlags;         // open(2) flags for (re)opening the file
    int fileMode;          // mode to pass to open(2)
} vfd;
```

Virtual file descriptors (**vfd**)

- physically stored in dynamically-allocated array



- also arranged into list by recency-of-use



**vfdCache[0]** holds list head/tail pointers.

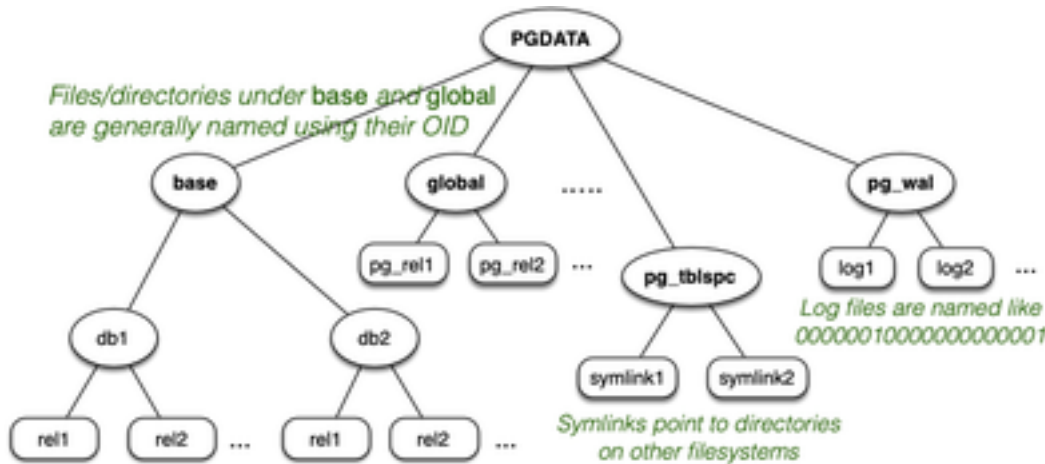
Consider the following call to open a file

```
f = PathNameOpenFilePerm(
    "/srvr/jas/pgsql/data/base/13645/12348",
    O_RDWR | O_CREAT | O_EXCL | PG_BINARY,
    0600
)
```

Sketch implementation of `PathNameOpenFilePerm()`

# File Manager

Reminder: PostgreSQL file organisation

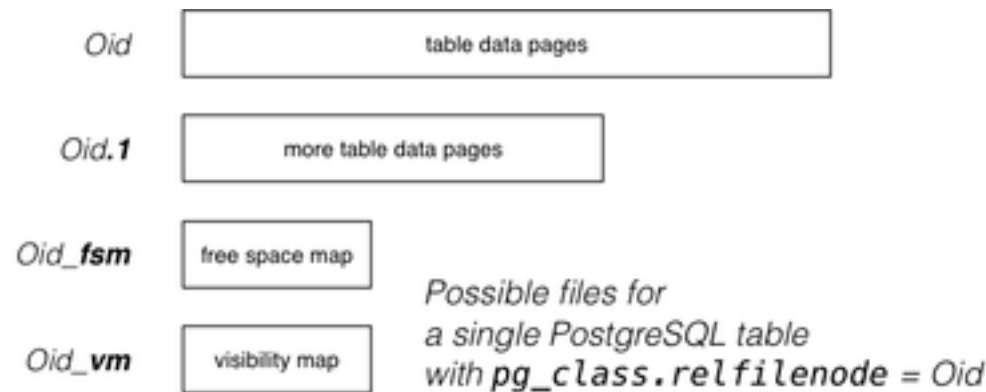


## ... File Manager

38/100

PostgreSQL stores each table

- in the directory *PGDATA/pg\_database.oid*
- often in multiple files (aka *forks*)




## ... File Manager

39/100

Data files (*Oid*, *Oid.1*, ...):

- sequence of fixed-size blocks/pages (typically 8KB)
- each page contains tuple data and admin data (see later)
- max size of data files 1GB (Unix limitation)



## ... File Manager

40/100

Free space map (*Oid\_fsm*):

- indicates where free space is in data pages
- "free" space is only free after **VACUUM**  
(**DELETE** simply marks tuples as no longer in use *xmax*)

Visibility map (*Oid\_vm*):

- indicates pages where all tuples are "visible"  
(*visible* = accessible to all currently active transactions)
- such pages can be ignored by **VACUUM**

## ... File Manager

41/100

The "magnetic disk storage manager" (**storage/smgr/md.c**)

- manages its own pool of open file descriptors (Vfd's)
- may use several Vfd's to access data, if several forks
- manages mapping from **PageID** to file+offset.

PostgreSQL PageID values are structured:

```
typedef struct
{
    RelFileNode rnode;    // which relation/file
    ForkNumber  forkNum;  // which fork (of reln)
    BlockNumber blockNum; // which page/block
} BufferTag;
```

---

## ... File Manager

42/100

Access to a block of data proceeds (roughly) as follows:

```
// pageID set from pg_catalog tables
// buffer obtained from Buffer pool
getBlock(BufferTag pageID, Buffer buf)
{
    File fid;  off_t offset;  int fd;
    (fid, offset) = findBlock(pageID)
    fd = VfdCache[fid].fd;
    lseek(fd, offset, SEEK_SET)
    VfdCache[fid].seekPos = offset;
    nread = read(fd, buf, BLOCKSIZE)
    if (nread < BLOCKSIZE) ... we have a problem
}
```

BLOCKSIZE is a global configurable constant (default: 8192)

---

## ... File Manager

43/100

```
findBlock(BufferTag pageID) returns (Vfd, off_t)
{
    offset = pageID.blockNum * BLOCKSIZE
    fileName = relpath(pageID.rnode)
    if (pageID.forkNum > 0)
        fileName = fileName+"."+pageID.forkNum
    fid = PathNameOpenFile(fileName, O_READ);
    fSize = VfdCache[fid].fileSize;
    if (offset > fSize) {
        fid = allocate new Vfd for next fork
        offset = offset - fd.fileSize
    }
    return (fd, offset)
}
```

---

# Buffer Pool

## Buffer Pool

45/100

Aim of buffer pool:

- hold pages read from database files, for possible re-use

Used by:

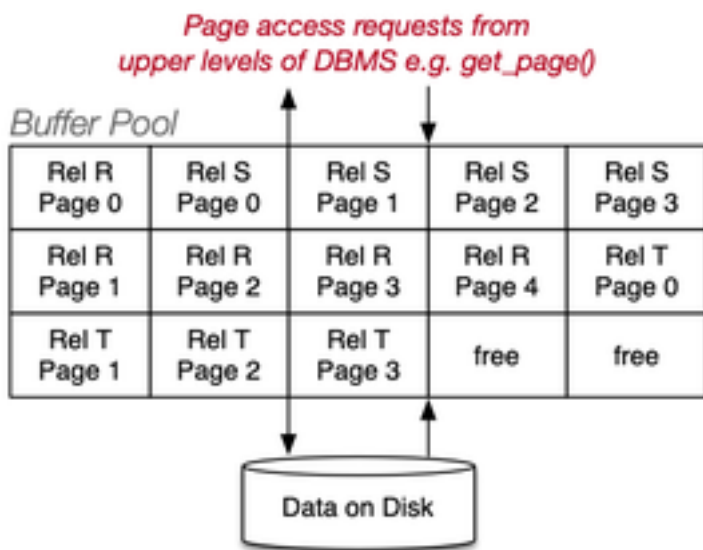
- *access methods* which read/write data pages
- e.g. sequential scan, indexed retrieval, hashing

Uses:

- file manager functions to access data files

Note: we use the terms *page* and *block* interchangeably

---



Buffer pool operations: (both take single `PageID` argument)

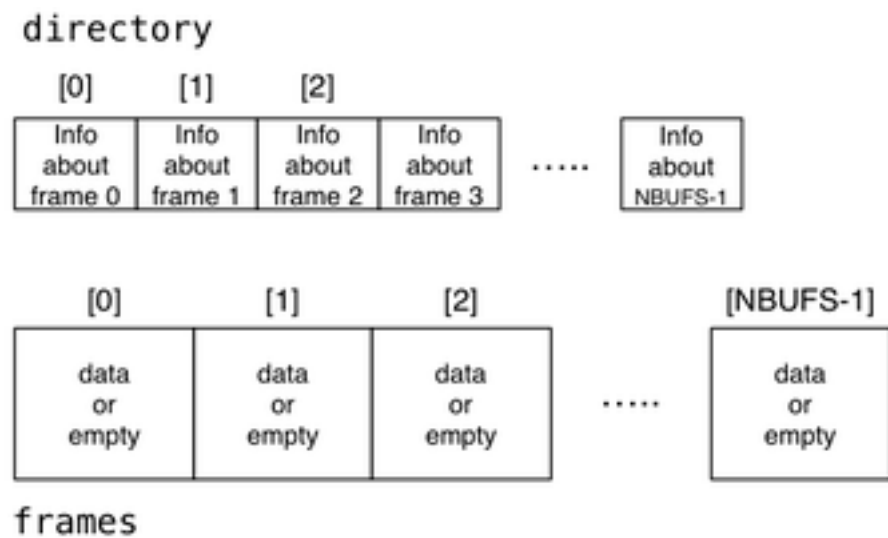
- `request_page(pid)`, `release_page(pid)`, ...

To some extent ...

- `request_page()` replaces `getBlock()` or `get_page()`
- `release_page()` replaces `putBlock()` or `put_page()`

Buffer pool data structures:

- `frames` ... array of `NBUFS` Page buffers
- `directory` ... array of `NBUFS` `FrameData` items



For each frame, we need to know: (`FrameData`)

- which Page it contains, or whether empty/free
- whether it has been modified since loading (*dirty bit*)
- how many transactions are currently using it (*pin count*)
- time-stamp for most recent access (assists with replacement)

Pages are referenced by `PageID` ...

- `PageID = BufferTag = (rnode, forkNum, blockNum)`

How scans are performed without Buffer Pool:

```

Buffer buf;
int N = numberOfBlocks(Rel);
for (i = 0; i < N; i++) {
    pageID = makePageID(db,Rel,i);
    getBlock(pageID, buf);
    for (j = 0; j < nTuples(buf); j++)
        process(buf, j)
}

```

Requires N page reads.

If we read it again, N page reads.

---

### ... Buffer Pool

51/100

How scans are performed with Buffer Pool:

```

Buffer buf;
int N = numberOfBlocks(Rel);
for (i = 0; i < N; i++) {
    pageID = makePageID(db,Rel,i);
    bufID = request_page(pageID);
    buf = frames[bufID]
    for (j = 0; j < nTuples(buf); j++)
        process(buf, j)
    release_page(pageID);
}

```

Requires N page reads on the first pass.

If we read it again,  $0 \leq \text{page reads} \leq N$

---

### ... Buffer Pool

52/100

Buffer pool data structures:

```

typedef char Page[PAGESIZE];
typedef ... PageID; // defined earlier

typedef struct _FrameData {
    PageID pid;           // which page is in frame
    int pin_count;        // how many processes using page
    int dirty;            // page modified since loaded?
    Time last_used;       // when page was last accessed
} FrameData;

Page frames[NBUFS];      // actual buffers
FrameData directory[NBUFS];

```

---

### ... Buffer Pool

53/100

Implementation of request\_page()

```

int request_page(PageID pid)
{
    bufID = findInPool(pid)
    if (bufID == NOT_FOUND) {
        if (no free frames in Pool) {
            bufID = findFrameToReplace()
            if (directory[bufID].dirty)
                old = directory[bufID].page
                put_page(old, frames[bufID])
        }
        bufID = index of freed frame
        directory[bufID].page = pid
        directory[bufID].pin_count = 0
        directory[bufID].dirty = 0
        get_page(pid, frames[bufID])
    }
    directory[bufID].pin_count++
    return bufID
}

```

The `release_page(pid)` operation:

- Decrement pin count for specified page

Note: no effect on disk or buffer contents until replacement required

The `mark_page(pid)` operation:

- Set dirty bit on for specified page

Note: doesn't actually write to disk; indicates that page changed

The `flush_page(pid)` operation:

- Write the specified page to disk (using `write()`)

Note: not generally used by higher levels of DBMS

---

Evicting a page ...

- find frame(s) preferably satisfying
  - pin count = 0 (i.e. nobody using it)
  - dirty bit = 0 (not modified)
- if selected frame was modified, flush frame to disk
- flag directory entry as "frame empty"

If multiple frames can potentially be released

- need a policy to decide which is best choice
- 

## Page Replacement Policies

Several schemes are commonly in use:

- Least Recently Used (LRU)
- Most Recently Used (MRU)
- First in First Out (FIFO)
- Random

LRU / MRU require knowledge of when pages were last accessed

- how to keep track of "last access" time?
  - base on request/release ops or on *real* page usage?
- 

Cost benefit from buffer pool (with  $n$  frames) is determined by:

- number of available frames (more  $\Rightarrow$  better)
- replacement strategy vs page access pattern

**Example (a):** sequential scan, LRU or MRU,  $n \geq b$

First scan costs  $b$  reads; subsequent scans are "free".

**Example (b):** sequential scan, MRU,  $n < b$

First scan costs  $b$  reads; subsequent scans cost  $b - n$  reads.

**Example (c):** sequential scan, LRU,  $n < b$

All scans cost  $b$  reads; known as *sequential flooding*.

---

## Effect of Buffer Management

Consider a query to find customers who are also employees:

```
select c.name
from Customer c, Employee e
where c.ssn = e.ssn;
```

This might be implemented inside the DBMS via nested loops:

```
for each tuple t1 in Customer {
    for each tuple t2 in Employee {
        if (t1.ssn == t2.ssn)
            append (t1.name) to result set
    }
}
```

... Effect of Buffer Management

59/100

In terms of page-level operations, the algorithm looks like:

```
Rel rC = openRelation("Customer");
Rel rE = openRelation("Employee");
for (int i = 0; i < nPages(rC); i++) {
    PageID pid1 = makePageID(db,rC,i);
    Page p1 = request_page(pid1);
    for (int j = 0; j < nPages(rE); j++) {
        PageID pid2 = makePageID(db,rE,j);
        Page p2 = request_page(pid2);
        // compare all pairs of tuples from p1,p2
        // construct solution set from matching pairs
        release_page(pid2);
    }
    release_page(pid1);
}
```

Exercise 4: Buffer Cost Benefit (i)

60/100

Assume that:

- the Customer relation has  $b_C$  pages (e.g. 10)
- the Employee relation has  $b_E$  pages (e.g. 4)

Compute how many page reads occur ...

- if we have only 2 buffers (i.e. effectively no buffer pool)
- if we have 20 buffers
- when a buffer pool with MRU replacement strategy is used
- when a buffer pool with LRU replacement strategy is used

For the last two, buffer pool has  $n=3$  slots ( $n < b_C$  and  $n < b_E$ )

Exercise 5: Buffer Cost Benefit (ii)

61/100

If the tables were larger, the above analysis would be tedious.

Write a C program to simulate buffer pool usage

- assuming a nested loop join as above
- `argv[1]` gives number of pages in "outer" table
- `argv[2]` gives number of pages in "inner" table
- `argv[3]` gives number of slots in buffer pool
- `argv[4]` gives replacement strategy (LRU,MRU,FIFO-Q)

PostgreSQL Buffer Manager

62/100

PostgreSQL buffer manager:

- provides a shared pool of memory buffers for all backends
- all access methods get data from disk via buffer manager

Buffers are located in a large region of *shared memory*.

Definitions: `src/include/storage/buf*.h`



Buffer code is also used by backends who want a private buffer pool

### ... PostgreSQL Buffer Manager

63/100

Buffer pool consists of:

**BufferDescriptors** (i.e. directory)

- shared fixed array of NBuffers x **BufferDesc**

**BufferBlocks** (i.e. frames)

- shared fixed array of NBuffers x buffers (each BLCKSZ bytes)

**Buffer** = index in above arrays

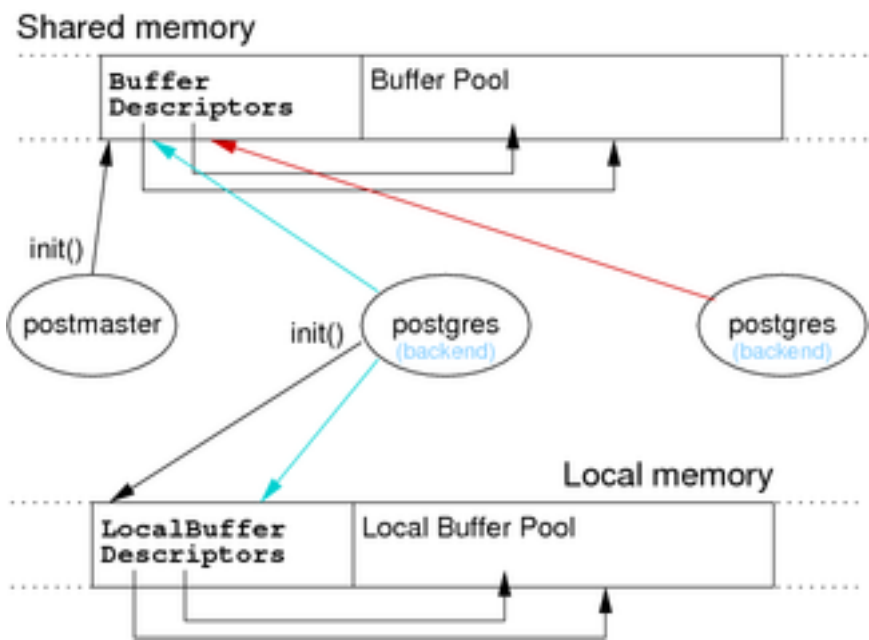
- indexes: global buffers 1..NBuffers; local buffers negative

Size of buffer pool is set in *postgresql.conf*, e.g.

shared\_buffers = 16MB # min 128KB, 16\*8KB buffers

### ... PostgreSQL Buffer Manager

64/100



## Buffer Pool Data Types

65/100

```
typedef struct buftag {
    RelFileNode rnode; /* physical relation identifier */
    ForkNumber forkNum;
    BlockNumber blockNum; /* relative to start of reln */
} BufferTag;

typedef struct BufferDesc { (simplified)
    BufferTag tag; /* ID of page contained in buffer */
    int buf_id; /* buffer's index number (from 0) */
    Bits32 state; /* dirty, refcount, usage */
    int freeNext; /* link in freelist chain */
    ... /* others related to concurrency */
} BufferDesc;
```

## Buffer Pool Functions

66/100

Buffer manager interface:

**Buffer ReadBuffer(Relation r, BlockNumber n)**

- ensures page  $n$  of file for relation  $r$  is loaded
- increments reference (pin) count and usage count for buffer
- returns index of loaded page in buffer pool (Buffer value)

```
BufferDesc *BufferAlloc(
    Relation r, ForkNumber f,
    BlockNumber n, bool *found)
```

- used by **ReadBuffer** to find a buffer for  $(r,f,n)$
- if no available buffers, select buffer to be replaced

---

## ... Buffer Pool Functions

67/100

Buffer manager interface (cont):

```
void ReleaseBuffer(Buffer buf)
```

- decrement pin count on buffer
- if pin count falls to zero, ensures all activity on buffer is completed before returning

```
void MarkBufferDirty(Buffer buf)
```

- marks a buffer as modified
- requires that buffer is pinned and locked
- actual write is done later (e.g. when buffer replaced)

---

## ... Buffer Pool Functions

68/100

Additional buffer manager functions:

```
Page BufferGetPage(Buffer buf)
```

- finds actual data associated with buffer in pool
- returns reference to memory where data is located

```
BufferIsPinned(Buffer buf)
```

- check whether this backend holds a pin on buffer

```
CheckpointBuffers
```

- write data in checkpoint logs (for recovery)
- flush all dirty blocks in buffer pool to disk

etc. etc. etc.

---

# Clock-sweep Replacement Strategy

69/100

PostgreSQL page replacement strategy: *clock-sweep*

- treat buffer pool as circular list of buffer slots
- NextVictimBuffer holds index of next possible evictee
- if this page is pinned or "popular", leave it
  - usage\_count implements "popularity/recency" measure
  - incremented on each access to buffer (up to small limit)
  - decremented each time considered for eviction
- increment NextVictimBuffer and try again (wrap at end)

---

# Exercise 6: PostgreSQL Buffer Pool

70/100

Consider an initially empty buffer pool with only 3 slots.

Show the state of the pool after each of the following:

```
Req R0, Req S0, Rel S0, Req S1, Rel S1, Req S2,
Rel S2, Rel R0, Req R1, Req S0, Rel S0, Req S1,
Rel S1, Req S2, Rel S2, Rel R1, Req R2, Req S0,
Rel S0, Req S1, Rel S1, Req S2, Rel S2, Rel R2
```

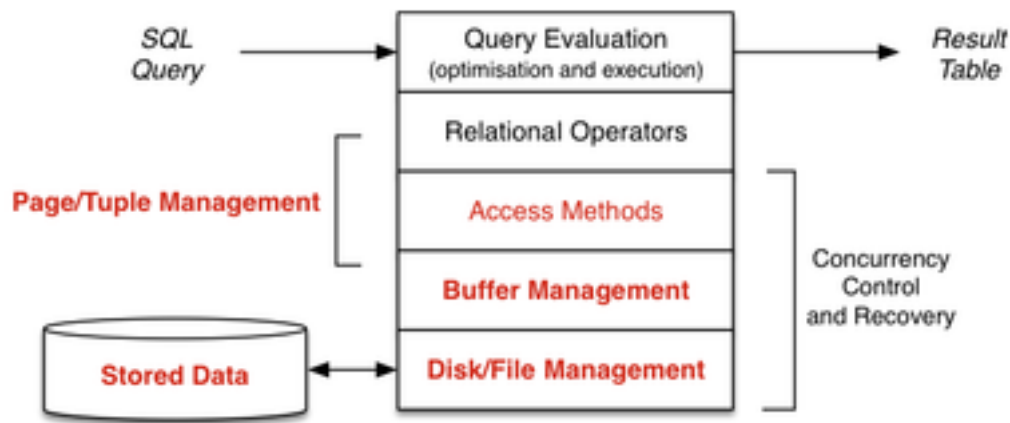
Treat BufferDesc entries as

Assume freeList and nextVictim global variables.

# Pages

## Page/Tuple Management

72/100



# Pages

73/100

Database applications view data as:

- a collection of records (tuples)
- records can be accessed via a TupleId/RecordId/RID
- $\text{TupleId} = (\text{PageID} + \text{TupIndex})$

The disk and buffer manager provide the following view:

- data is a sequence of fixed-size pages (aka "blocks")
- pages can be (random) accessed via a PageID
- each page contains zero or more tuple values

Page format = how space/tuples are organised within a page

## Page Formats

74/100

Ultimately, a Page is simply an array of bytes (byte[ ]).

We want to interpret/manipulate it as a collection of Records.

Typical operations on Pages:

- request\_page(pid) ... get page via its PageId
- get\_record(rid) ... get record via its TupleId
- rid = insert\_record(pid,rec) ... add new record
- update\_record(rid,rec) ... update value of record
- delete\_record(rid) ... remove record from page

Note: rid contains (PageId,TupIndex), so no explicit pid needed

### ... Page Formats

75/100


Factors affecting Page formats:

- determined by record size flexibility (fixed, variable)
- how free space within Page is managed
- whether some data is stored outside Page
  - does Page have an associated overflow chain?
  - are large data values stored elsewhere? (e.g. TOAST)
  - can one tuple span multiple Pages?

Implementation of Page operations critically depends on format.

For fixed-length records, use *record slots*.

- *insert*: place new record in first available slot
- *delete*: two possibilities for handling free record slots:



Exercise 7: Fixed-length Records

77/100

Give examples of table definitions

- which result in fixed-length records
- which result in variable-length records

```
create table R ( ...);
```

What are the common features of each type of table?

Exercise 8: Inserting/Deleting Fixed-length Records

78/100

For each of the following Page formats:

- compacted/packed free space
- unpacked free space (with bitmap)

Implement

- a suitable data structure to represent a Page
- a function to insert a new record
- a function to delete a record

Page Formats

79/100

For variable-length records, must use *slot directory*.

Possibilities for handling free-space within block:

- compacted (one region of free space)
- fragmented (distributed free space)


In practice, a combination is useful:

- normally fragmented (cheap to maintain)
- compacted when needed (e.g. record won't fit)

Important aspect of using slot directory

- location of tuple within page can change, tuple index does not change

Compacted free space:

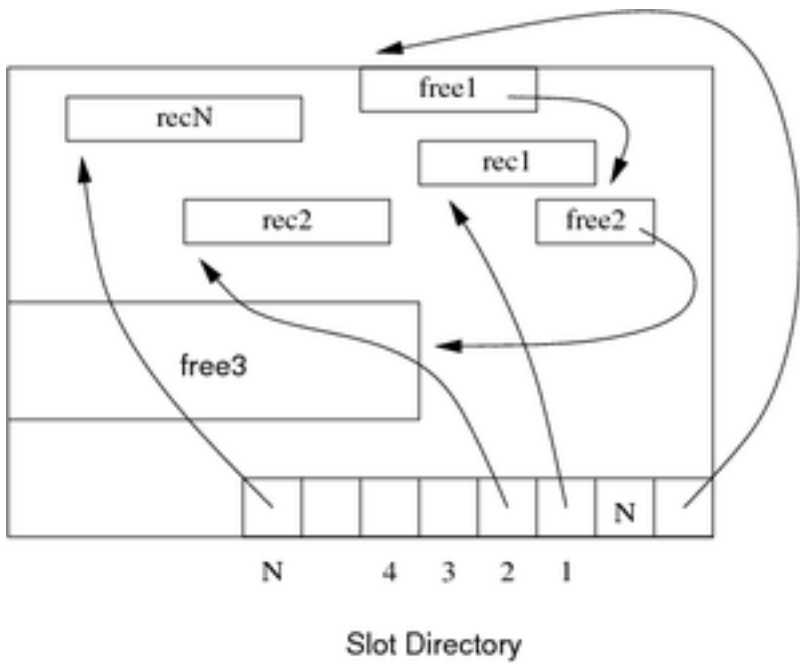


Note: "pointers" are implemented as word offsets within block.

... Page Formats

81/100


Fragmented free space:



... Page Formats

82/100


Initial page state (compacted free space) ...



... Page Formats

83/100


Before inserting record 7 (compacted free space) ...



... Page Formats

84/100


After inserting record 7 (80 bytes) ...



... Page Formats

85/100

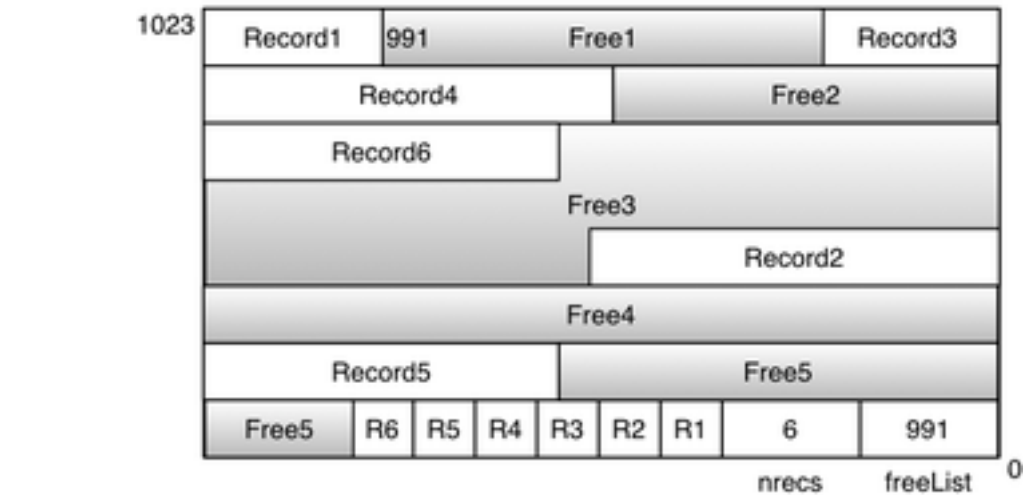
Initial page state (fragmented free space) ...



... Page Formats

86/100

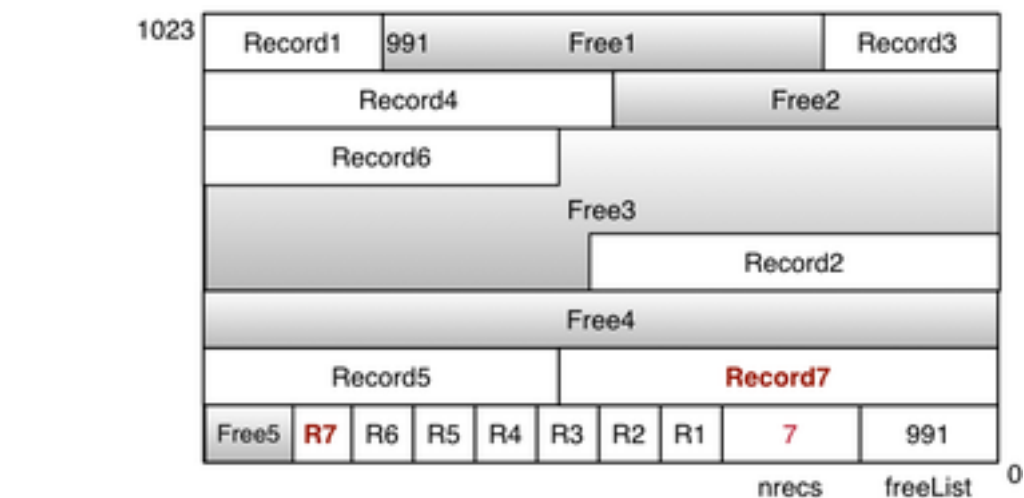
Before inserting record 7 (fragmented free space) ...



... Page Formats

87/100

After inserting record 7 (80 bytes) ...



Exercise 9: Inserting Variable-length Records

88/100

For both of the following page formats

1. variable-length records, with compacted free space
2. variable-length records, with fragmented free space

implement the `insert()` function.

Use the above page format, but also assume:

- page size is 1024 bytes
- tuples start on 4-byte boundaries
- references into page are all 8-bits (1 byte) long
- a function `recSize(r)` gives size in bytes

Storage Utilisation

89/100

How many records can fit in a page? (denoted  $C$  = capacity)

Depends on:

- page size ... typical values: 1KB, 2KB, 4KB, 8KB
- record size ... typical values: 64B, 200B, app-dependent
- page header data ... typically: 4B - 32B
- slot directory ... depends on how many records

We typically consider *average* record size ( $R$ )

Given  $C$ ,  $HeaderSize + C*SlotSize + C*R \leq PageSize$

Exercise 10: Space Utilisation

90/100

Consider the following page/record information:



- page size = 1KB = 1024 bytes =  $2^{10}$  bytes
- records: `(a:int,b:varchar(20),c:char(10),d:int)`
- records are all aligned on 4-byte boundaries
- `c` field padded to ensure `d` starts on 4-byte boundary
- each record has 4 field-offsets at start of record (each 1 byte)
- `char(10)` field rounded up to 12-bytes to preserve alignment
- maximum size of `b` values = 20 bytes; average size = 16 bytes
- page has 32-bytes of header information, starting at byte 0
- only insertions, no deletions or updates

Calculate  $C$  = average number of records per page.

## Overflows

91/100

Sometimes, it may not be possible to insert a record into a page:

1. no free-space fragment large enough
2. overall free-space is not large enough
3. the record is larger than the page
4. no more free directory slots in page

For case (1), can first try to compact free-space within the page.

If still insufficient space, we need an alternative solution ...

### ... Overflows

92/100

File organisation determines how cases (2)..(4) are handled.

If records may be inserted anywhere that there is free space

- cases (2) and (4) can be handled by making a new page
- case (3) requires either spanned records or "overflow file"

If file organisation determines record placement (e.g. hashed file)


- cases (2) and (4) require an "overflow page"
- case (3) requires an "overflow file"

With overflow pages, *rid* structure may need modifying (*rel*,*page*,*ovfl*,*rec*)


### ... Overflows

93/100

Overflow files for very large records and BLOBs:



Record-based handling of overflows:



We discuss overflow pages in more detail when covering Hash Files.

## PostgreSQL Page Representation

94/100

Functions: `src/backend/storage/page/*.c`

Each page is 8KB (default `BLCKSZ`) and contains:

- header (free space pointers, flags, xact data)
- array of (offset,length) pairs for tuples in page
- free space region (between array and tuple data)
- actual tuples themselves (inserted from end towards start)
- (optionally) region for special data (e.g. index data)

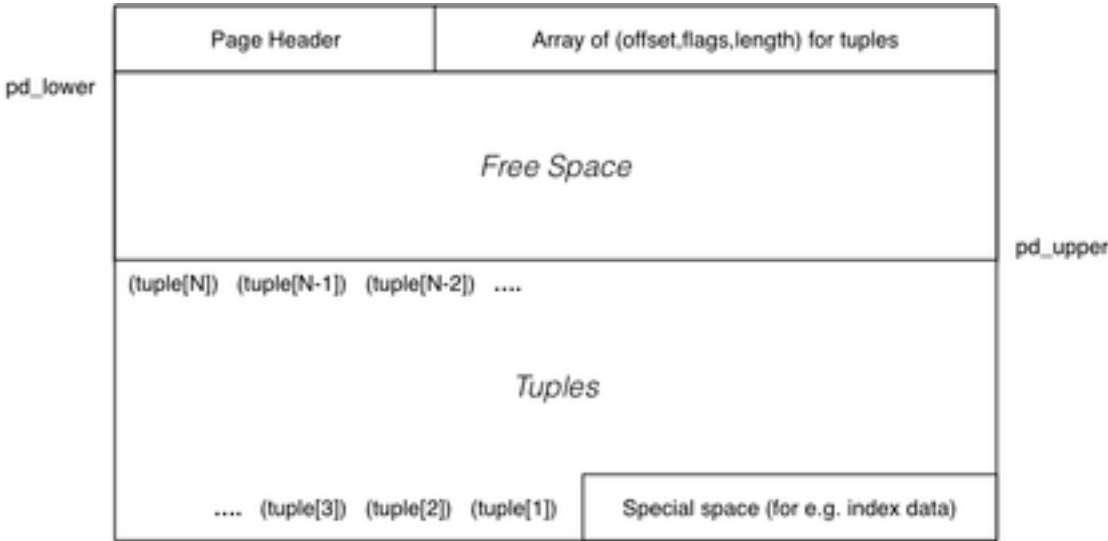
Large data items are stored in separate (TOAST) files (implicit)

Also supports ~SQL-standard BLOBs (explicit large data items)

... PostgreSQL Page Representation

95/100

PostgreSQL page layout:



... PostgreSQL Page Representation

96/100

Page-related data types:

```
// a Page is simply a pointer to start of buffer
typedef Pointer Page;

// indexes into the tuple directory
typedef uint16 LocationIndex;

// entries in tuple directory (line pointer array)
typedef struct ItemIdData
{
    unsigned    lp_off:15,      // tuple offset from start of page
                lp_flags:2,     // unused,normal,redirect,dead
                lp_len:15;      // length of tuple (bytes)
} ItemIdData;
```

... PostgreSQL Page Representation

97/100

Page-related data types: (cont)

```
typedef struct PageHeaderData
{
    XLogRecPtr    pd_lsn;        // xact log record for last change
    uint16        pd_tli;        // xact log reference information
    uint16        pd_flags;      // flag bits (e.g. free, full, ...)
    LocationIndex pd_lower;      // offset to start of free space
    LocationIndex pd_upper;      // offset to end of free space
    LocationIndex pd_special;    // offset to start of special space
    uint16        pd_pagesize;   // version;
    TransactionId pd_prune_xid;  // is pruning useful in data page?
    ItemIdData    pd_linp[1];    // beginning of line pointer array
} PageHeaderData;

typedef PageHeaderData *PageHeader;
```

... PostgreSQL Page Representation

98/100

Operations on Pages:

**void PageInit(Page page, Size pageSize, ...)**

- initialize a Page buffer to empty page
- in particular, sets pd\_lower and pd\_upper

**OffsetNumber PageAddItem(Page page,  
Item item, Size size, ...)**

- insert one tuple (or index entry) into a Page
- fails if: not enough free space, too many tuples

**void PageRepairFragmentation(Page page)**

- compact tuple storage to give one large free space region

---

## ... PostgreSQL Page Representation

99/100

PostgreSQL has two kinds of pages:

- *heap pages* which contain tuples
- *index pages* which contain index entries

Both kinds of page have the same page layout.

One important difference:

- index entries tend be a smaller than tuples
- can typically fit more index entries per page

---

## Exercise 11: PostgreSQL Pages

100/100

Draw diagrams of a PostgreSQL heap page

- when it is initially empty
- after three tuples have been inserted  
with lengths of 60, 80, and 70 bytes
- after the 80 byte tuple is deleted (but before vacuuming)
- after a new 50 byte tuple is added

Show the values in the tuple header.

Assume that there is no special space in the page.